

O'REILLY®

Hands-on Hugging Face for Natural Language Processing





Prerequisites

- Comfortable programming in Python
- Basic familiarity with machine learning and deep learning
- Basic familiarity with natural language processing



Poll Question

How comfortable are you building and training machine learning models in Python?

- Never worked with ML before
- Built traditional ML models using scikit-learn
- Built ML models using TensorFlow and PyTorch



Poll Question

Have you worked with Natural Language Processing before?

- Never worked with NLP before
- Built traditional NLP models with scikit-learn
- Built NLP models with TensorFlow, PyTorch



Hugging Face

A company known for its work in the field of artificial intelligence particularly in the field of natural language processing



Hugging Face

A machine learning and data science platform and community that helps users build, deploy, and train machine learning models



Natural Language Processing with Transformers

- The original transformer architecture was introduced in the 2017 paper “**Attention Is All You Need**”
- Introduced the novel concept of attention to focus on the right parts of the input
- Revolutionized the field of NLP, leading to the powerful LLMs that we work with today





Hugging Face Transformer Library

- Offer a wide range of **pre-trained** transformer models for different use cases
- Cross-framework compatibility – all models can be used with both PyTorch and TensorFlow
- Used in a variety of applications – chatbots, enhancing search, translation, and more



O'REILLY®

Attention and Transformer Models



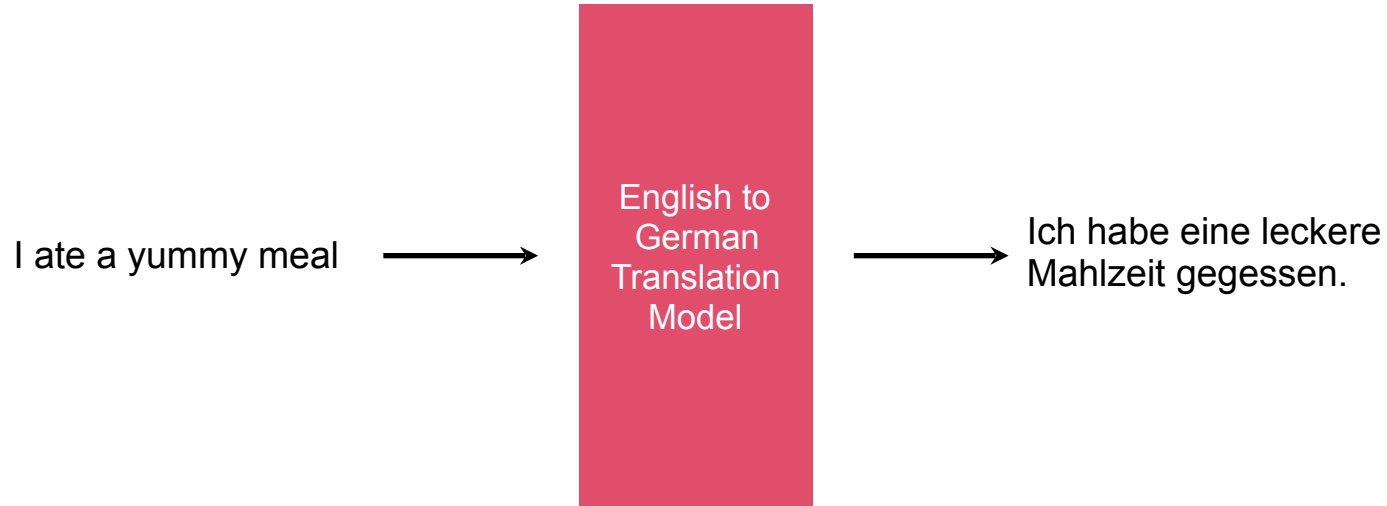


Attention in Neural Networks

Mechanism that allows the network to focus on specific parts of the input data while ignoring others



Language Translation Model





Model Without Attention

I ate a yummy meal



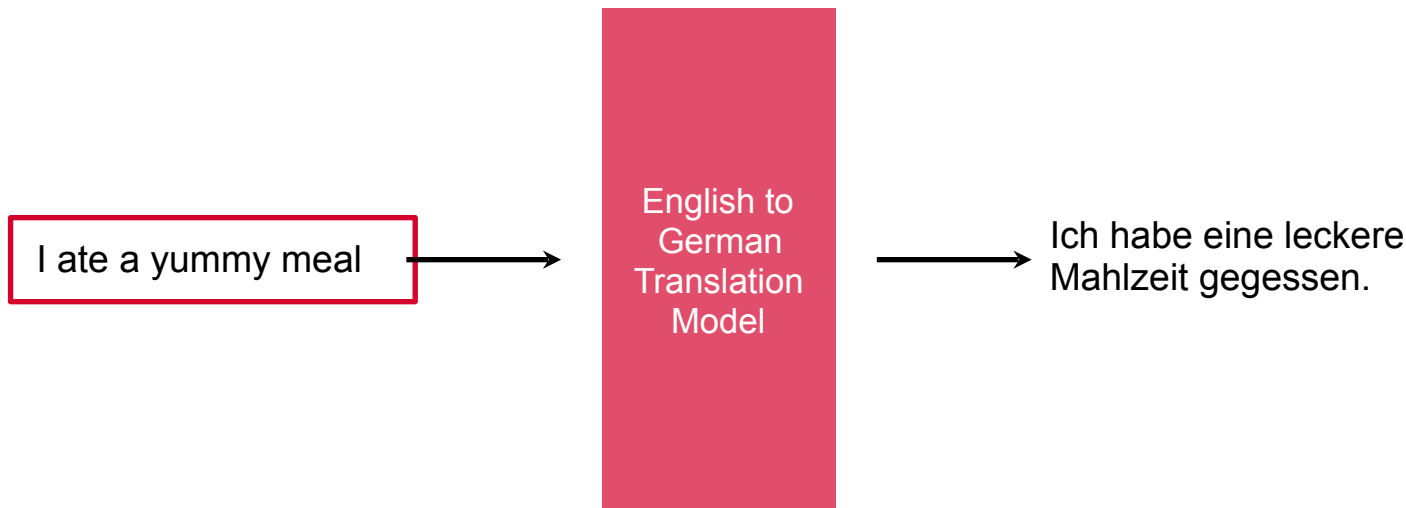
English to
German
Translation
Model



Ich habe eine leckere
Mahlzeit gegessen.

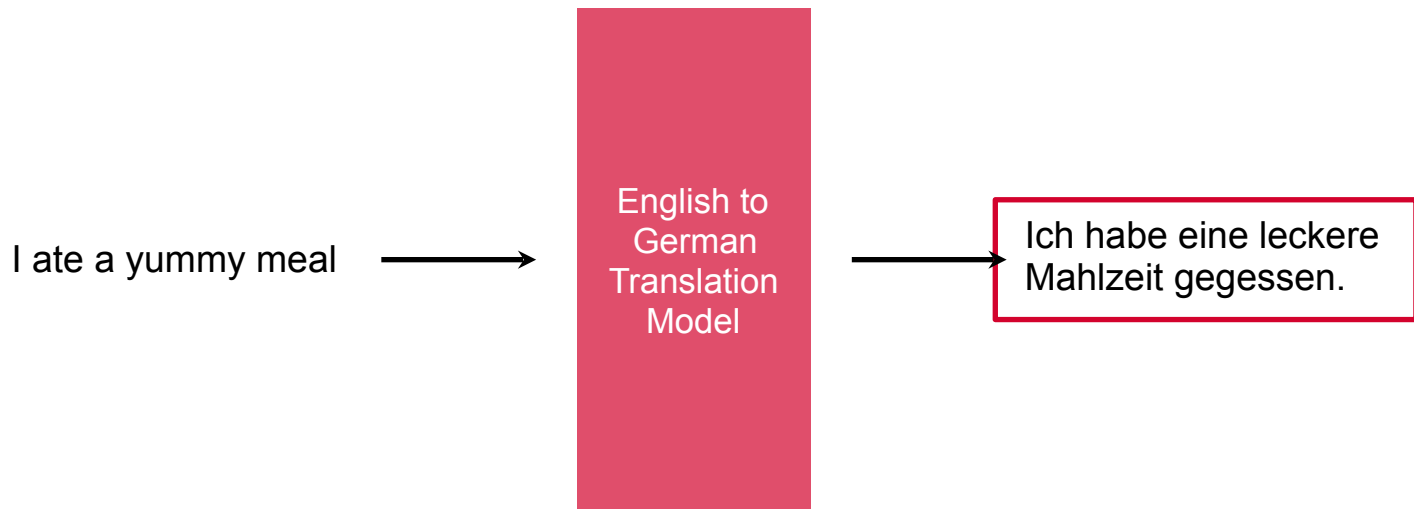


Parses the Entire Input






To Generate the Translation



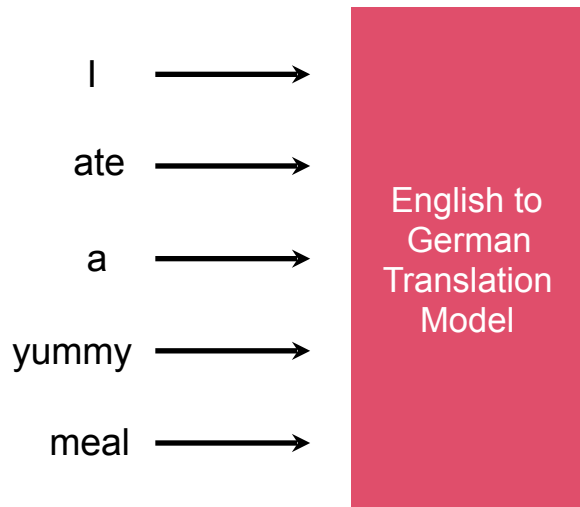


Input and Output are Sequences



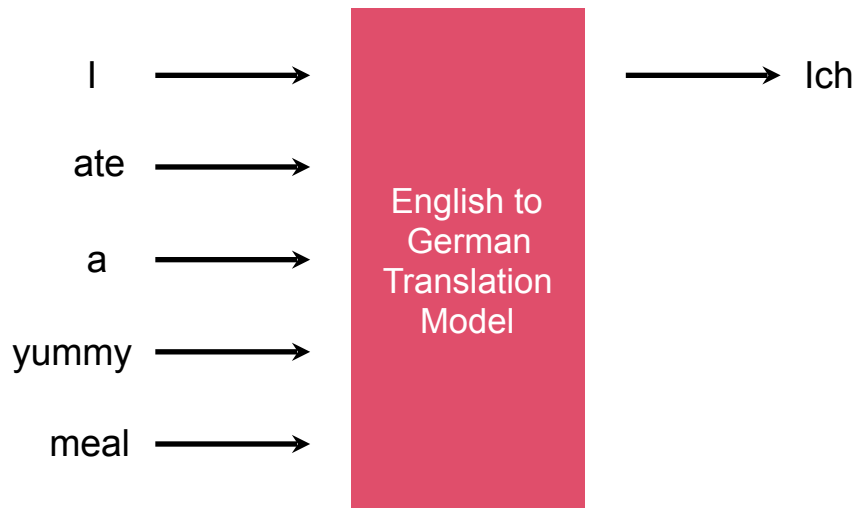


Input Sequence Fed In



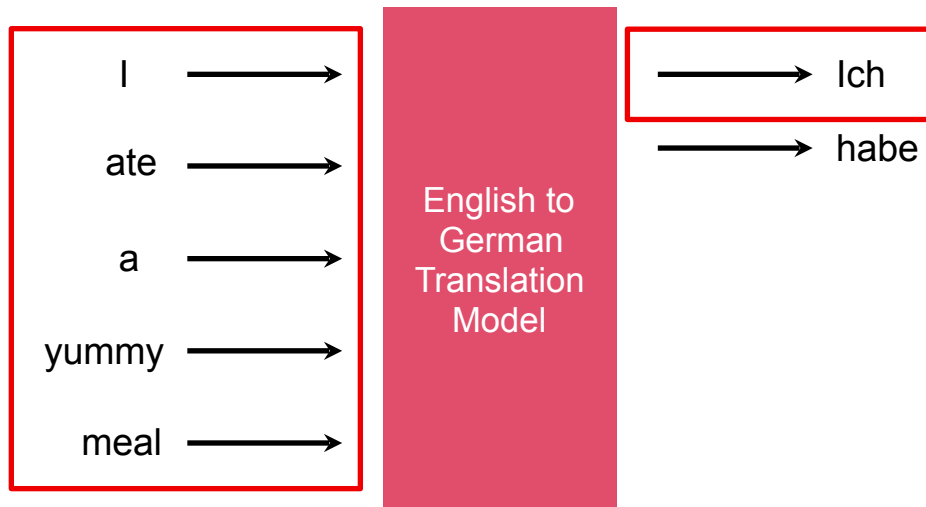


Output Sequence Produced one Word at a Time



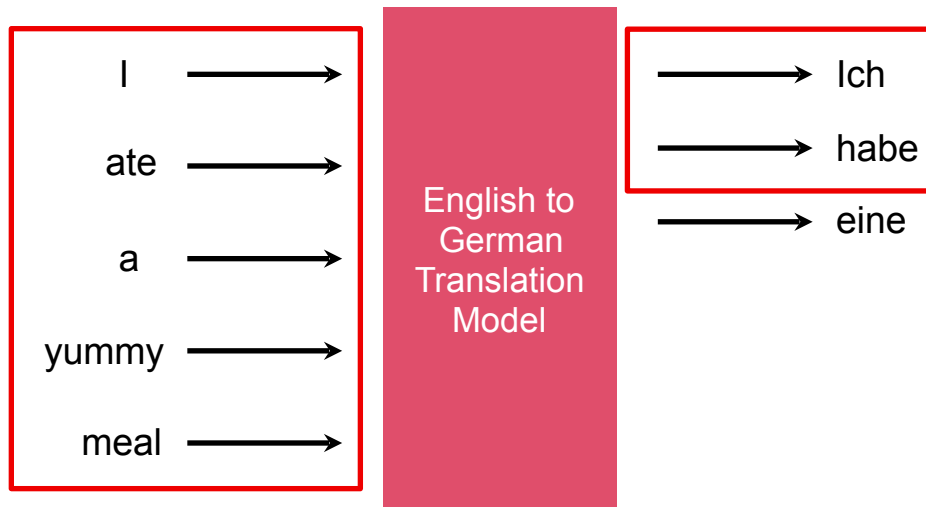


Complete Input + Words Produces So Far Used to Get Next Word in Sequence



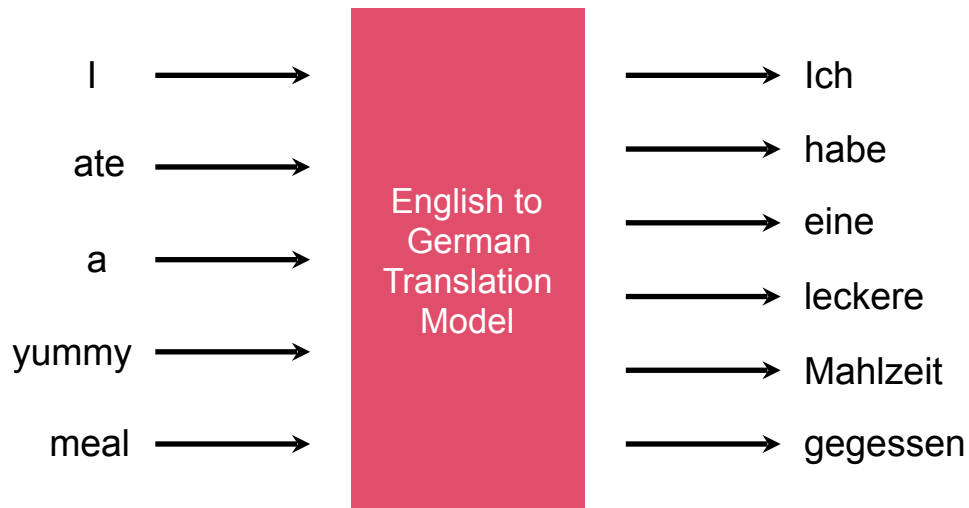


Predict One Word at a Time



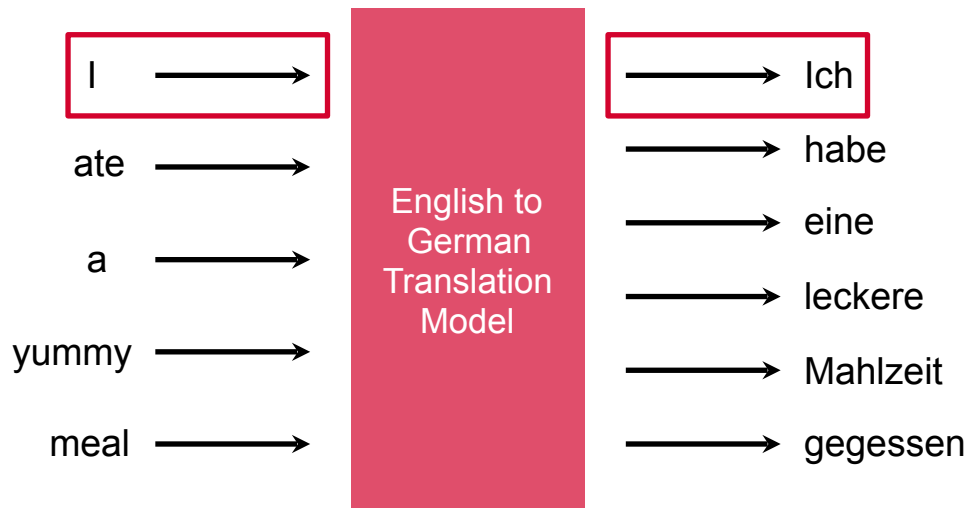


Better Predictions with Attention



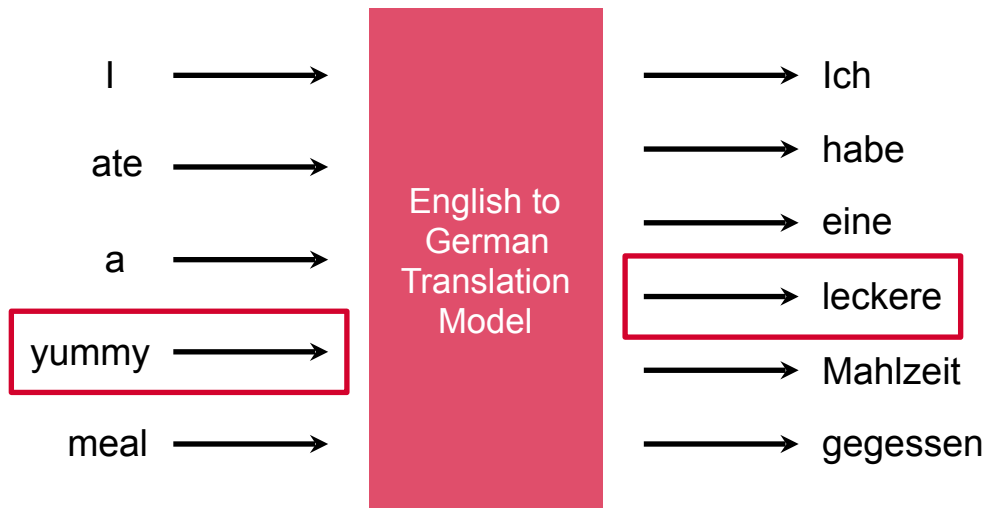


Better Predictions with Attention



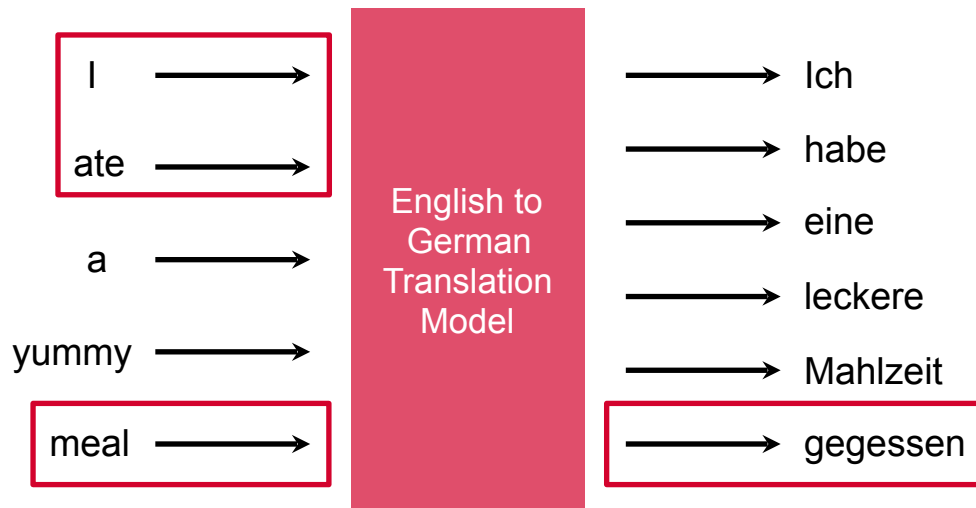


Better Predictions with Attention





Better Predictions with Attention





Attention

- Help models focus on the right parts of the input to generate the output
- The entire input along with attention generates better outputs
- Are the foundation of transformers used to power the large language models of today



O'REILLY®

The Transformer Architecture






The Transformer Architecture


A model architecture primarily used in natural language processing that relies on **self-attention** mechanisms to process sequential data more effectively than previous models

High-level Architecture of Transformers




Encoder Creates Abstract Representation

Fed to the Decoder

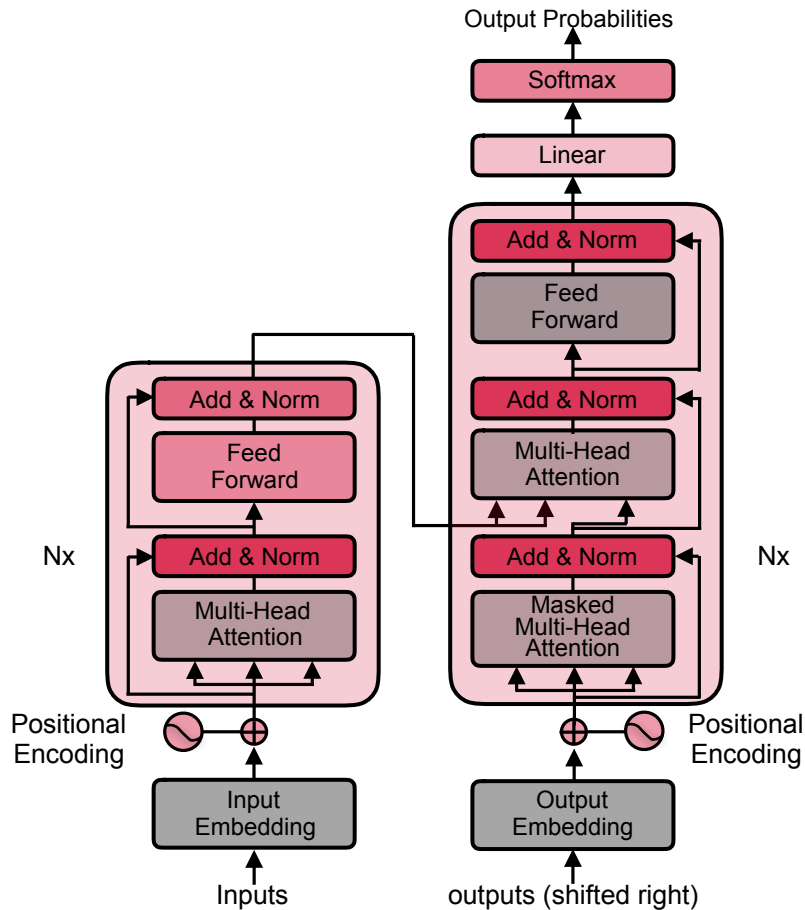


Encoder Creates Abstract Representation Fed to the Decoder

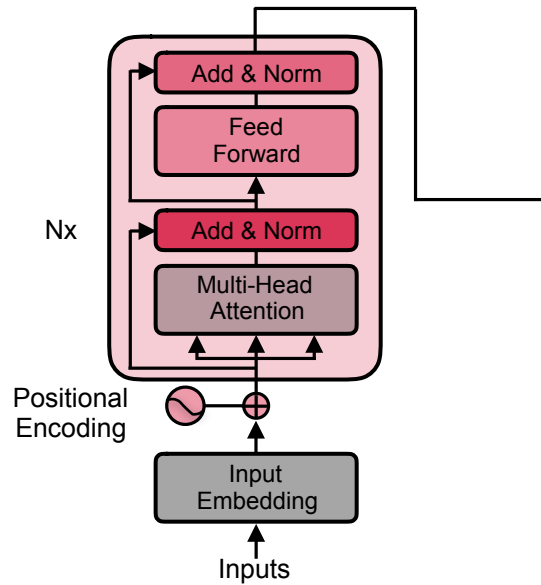


This representation encodes the importance of every word with respect to other words in the same input sequence

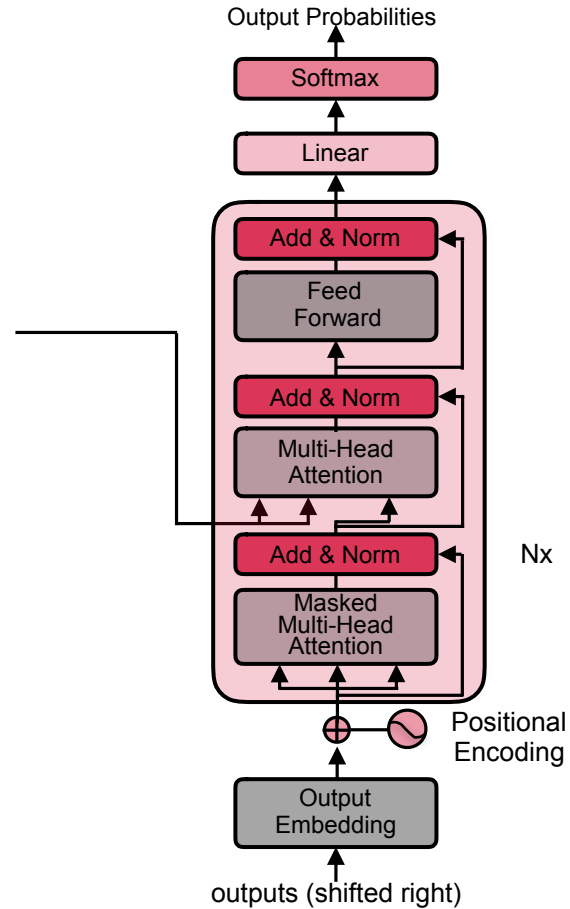
Transformer Architecture



Encoder



Decoder



O'REILLY®

Transformers in Hugging Face



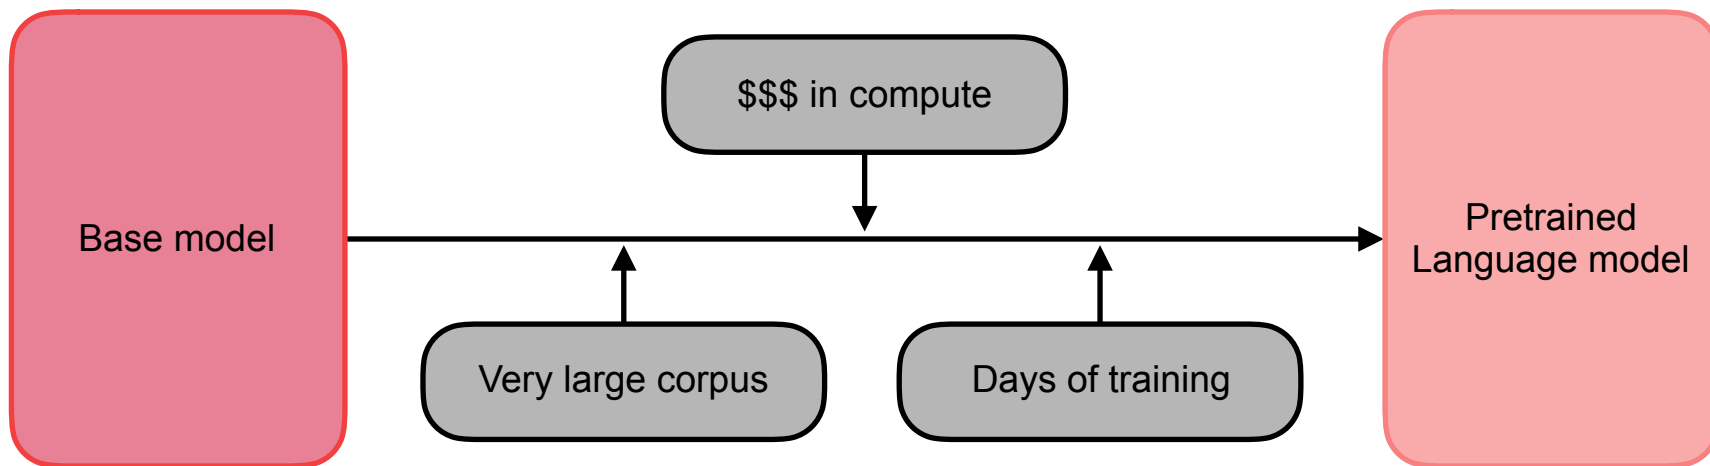


Pretrained Transformers in Hugging Face

Transformers are very larger models and very expensive to train – in terms of data, time, and resources




Training Transformers from Scratch



Training transformer models from scratch (pretraining) - model weights are initialized at random and the model has no prior knowledge embedded within it



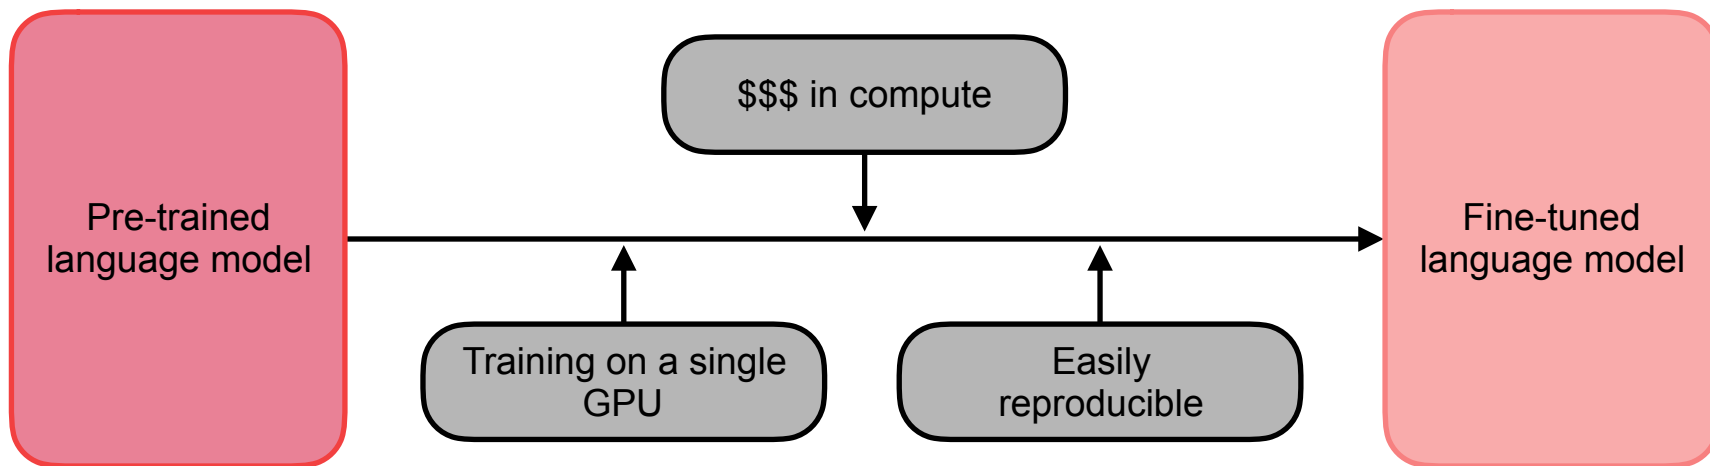
Training Transformers from Scratch



Needs huge amounts of data and compute resources - can be very expensive



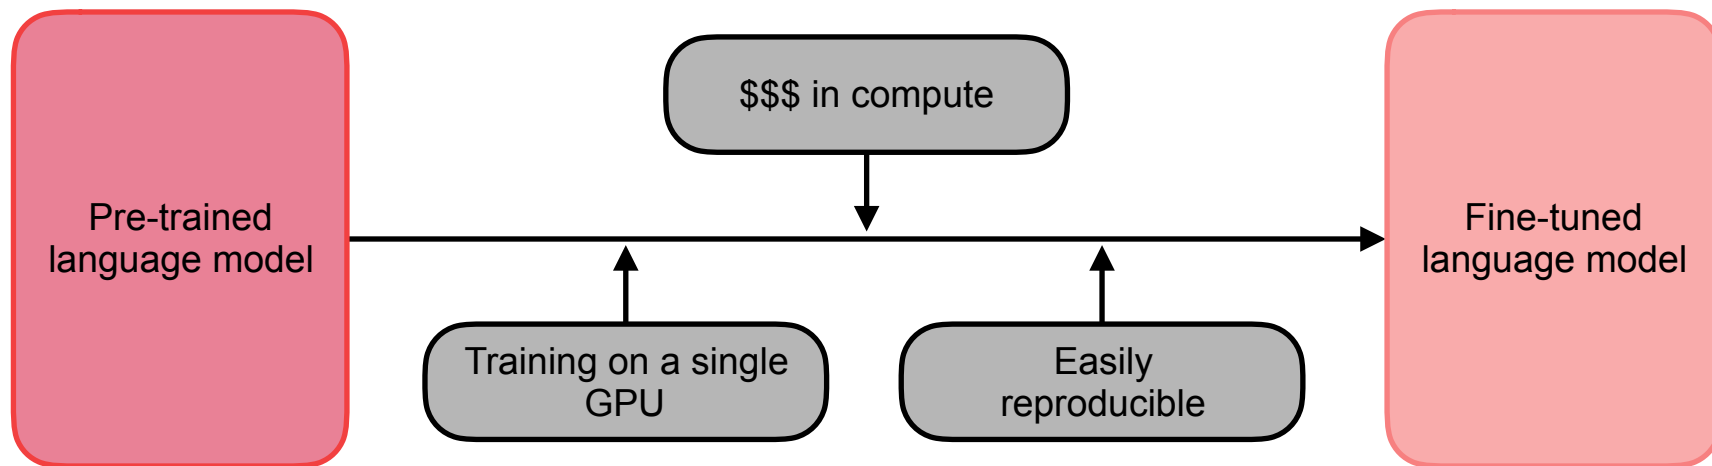
Fine-tuning Transformers



Fine-tuning models involve additional training on an already pretrained model



Fine-tuning Transformers



Use less data, time, and resources to get good results - overall costs and environment impact of training is much lower than training from scratch




Hugging Face Pipelines

A high-level abstraction that makes it easy for users to perform NLP tasks using pre-trained models



Pipeline Components




O'REILLY®

Hugging Face Tokenizers






Pipeline Components



Tokenizer





Tokenizing Text

Don't you know what time it is?



Split by Spaces

["Don't", "you", "know", "what", "time", "it", "is?"]



Consider Punctuation

["Don", "'", "t", "you", "know", "what", "time", "it", "is", "?"]



Split Words Like Won't and Don't

["Do", "n't", "you", "know", "what", "time", "it", "is", "?"]

Tokenizers

- Rule-based tokenizers with different rules split data in different ways
- Pretrained models only work well with **inputs tokenized in the same manner as its training data**





Categories of Tokenizers

Word Tokenizers

Character Tokenizers

Subword Tokenizers



Word and Character Tokenizers

- Split text into tokens in meaningful ways
- “swim” and “swimming” should be identified as related
- Reduce the number of unknown words
- Word tokenizers
 - Too large a vocabulary
 - Will not identify related words
- Character tokenizers
 - Large number of tokens for each sentence
 - Intuitively less meaningful





Subword Tokenization

- Relies on the principle that frequently used words should not be split but rare words should be split into subwords
- “emotionally” could be a rare word, split into “emotional” and “ly”
- Allow models to have a **reasonable vocabulary size and yet learn meaningful context independent representations**




O'REILLY®

Hugging Face Pipelines








Pipeline Components






Hugging Face Transformer Pipeline






Hugging Face Transformer Pipeline







General Transformer Architecture



Encoder



Decoder



Encoder-only Models

- Attention layers in the encoder can access all the words in the input sentence called “bi-directional” or “auto-encoding” models
- Pre-training involves corrupting the initial sentence (masking) and having the model predict the masked word
- Suited for tasks such as **sentence classification, named entity recognition, and question answering**





Masked Word Prediction

Let's have some fun doing _____ programming!

Python?

Java?

C++?



Decoder-only Models

- Attention layers in the decoder can only access words before the current word in the sentence – called “auto-regressive” models
- Pre-training of these models involve predicting the next word in the sentence, given all words seen so far
- Models are best suited for natural language tasks such as **text generation**





Next Word Prediction

It's a nice sunny afternoon. Let's go _____

outside

swimming

running

Dancing

Encoder-Decoder Models

- Called sequence-to-sequence models
- Attention layers in the encoder can access all parts of the input sentence,
- Attention in the decoder can only access words positioned before the given word
- Pre-training usually done using the objectives of the encoder and decoder (masked word prediction and next word prediction)
- Suited for tasks such as **summarization**, **translation**, and **generative question answering**

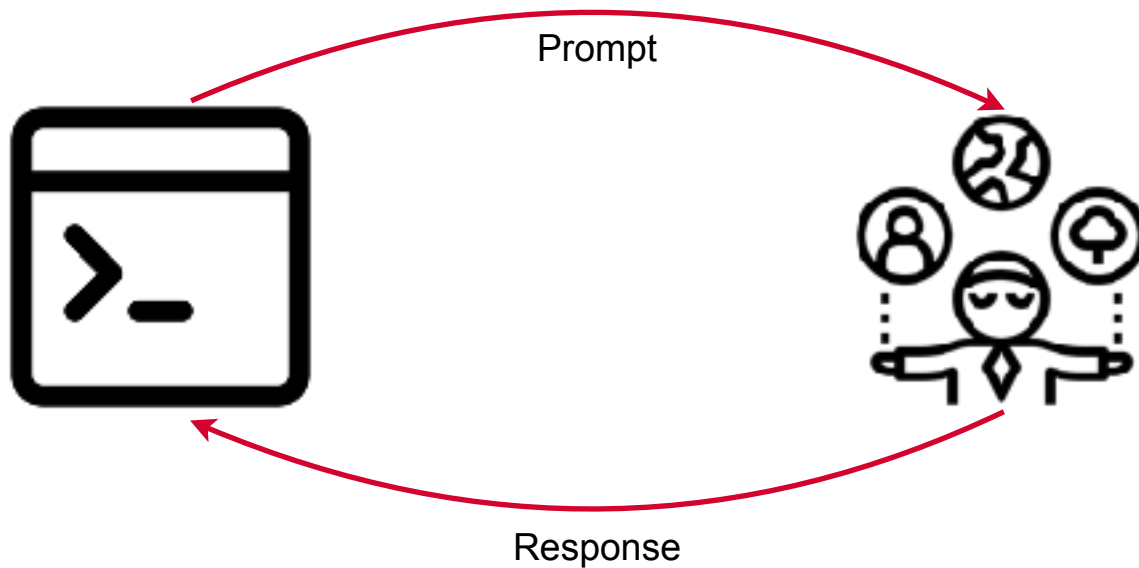


O'REILLY®

Text Generation by LLMs



Model Responses to Prompts





Response a Sequence of Words

She speaks French quite well




Model Generates One Word at a Time


She → speaks → French → quite → well




Many Possible Words at Each Step




Each Possible Next Word is Assigned a Probability



The Model Picks One Word from Possible Next Words



Higher Probability Words are More Likely to be Picked



Response Generated by Picking Words at Each Step



She → speaks → French → quite → well



Model Settings for Tweaking Generated Text

Large language models offer settings that you can tweak to make the generated text **more creative and diverse or more predictable and deterministic**



Creativity vs. Predictability in Text Generation

- High creativity will produce more diverse and unexpected results making the text more engaging
- High predictability generates more consistent and reliable text – useful when you need precise responses
- Striking a balance can produce text that is both interesting and coherent



Model Settings to Control Creativity and Predictability

- Temperature
- Top-p (Nucleus Sampling)
- Top-k



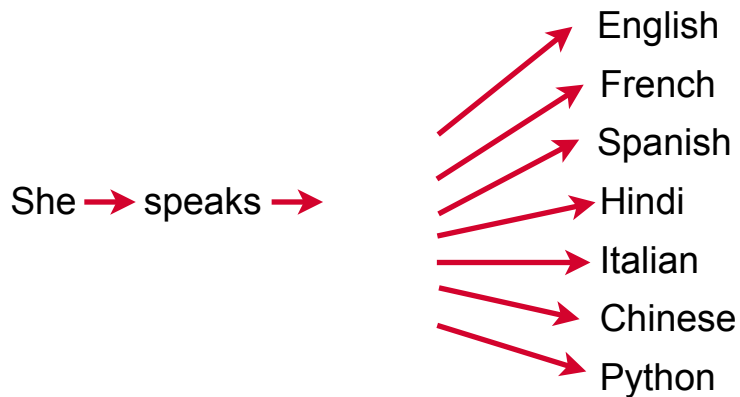
Temperature



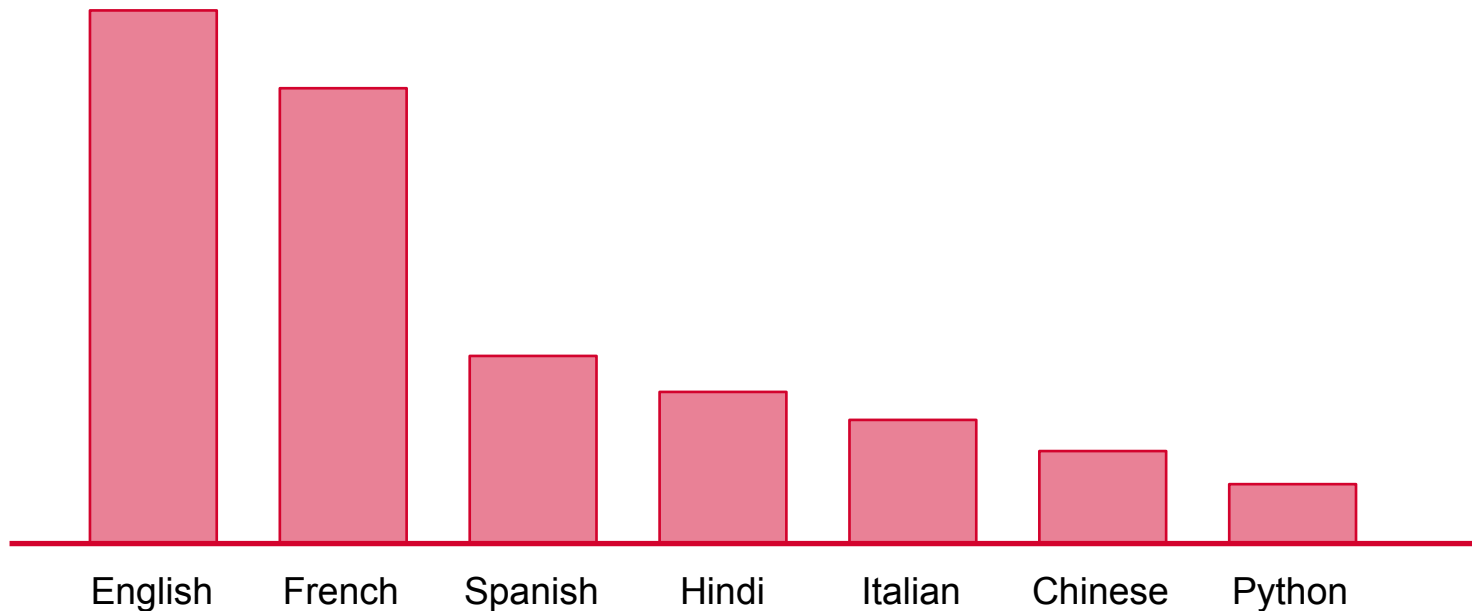
- Values range between 0 and 1 (both inclusive)
- Higher values closer to 1 results in more creative output
- Lower values closer to 0 results in more predictable output



The Model Picks One Word from Possible Next Words




Original Probabilities of Possible Next Words



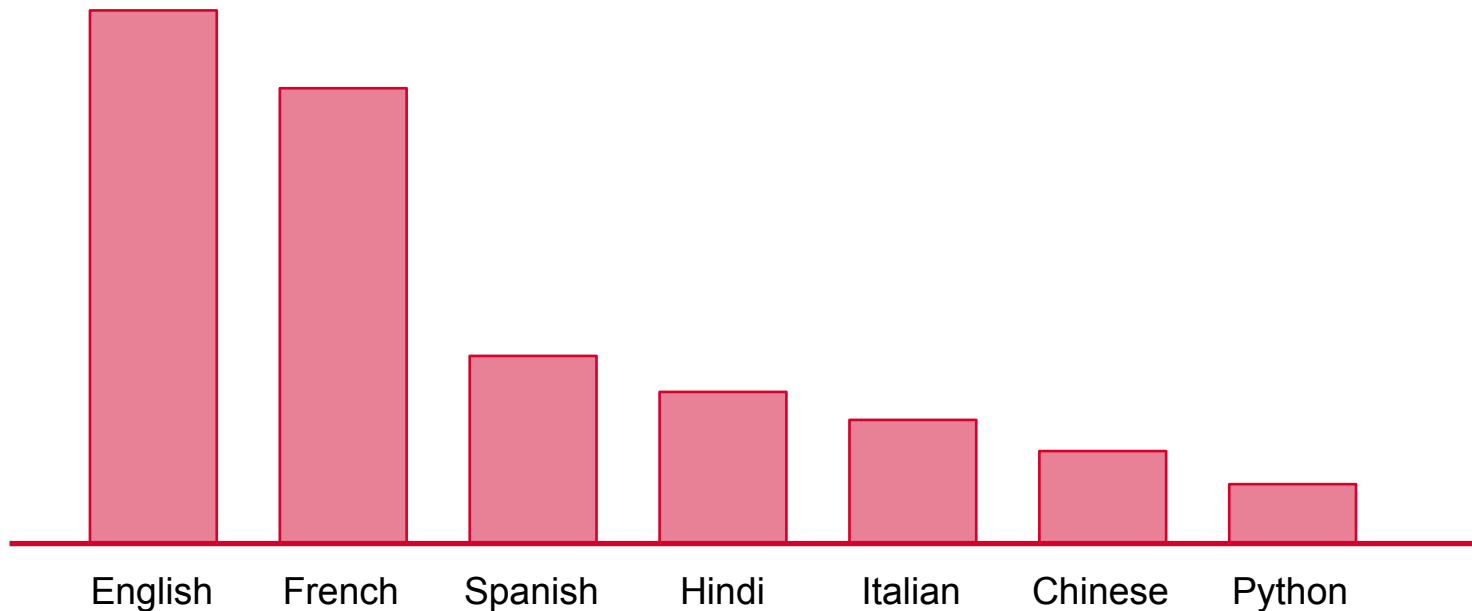


Higher Values of Temperature (Closer to 1)



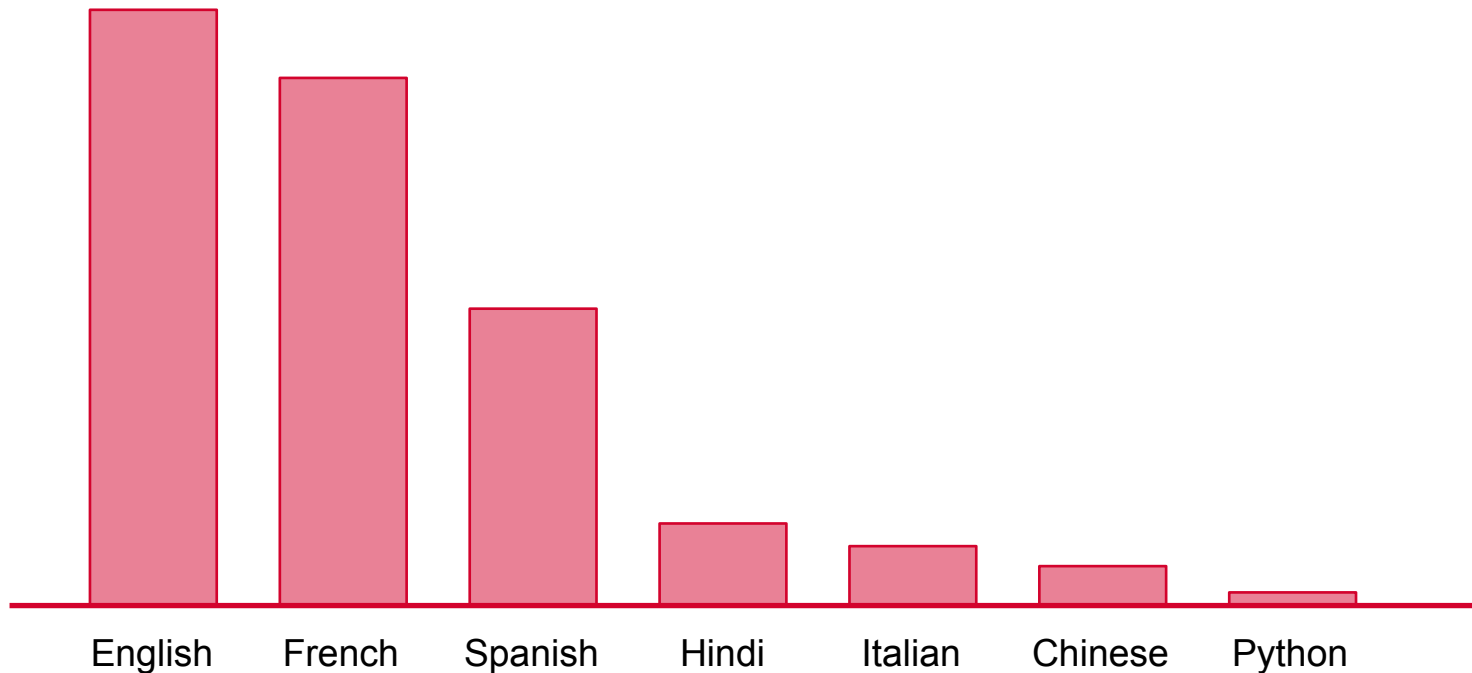
The probability distribution over the next possible word becomes **flatter**.

Original Probabilities of Possible Next Words





Lower Values of Temperature (Closer to 0)



The probability distribution over the next possible word becomes **sharper**.



Top-p (Nucleus Sampling)

- Values range between 0 and 1 (both inclusive)
- Values close to 1 result in more diverse and creative output
- Values close to 0 result in more predictable output

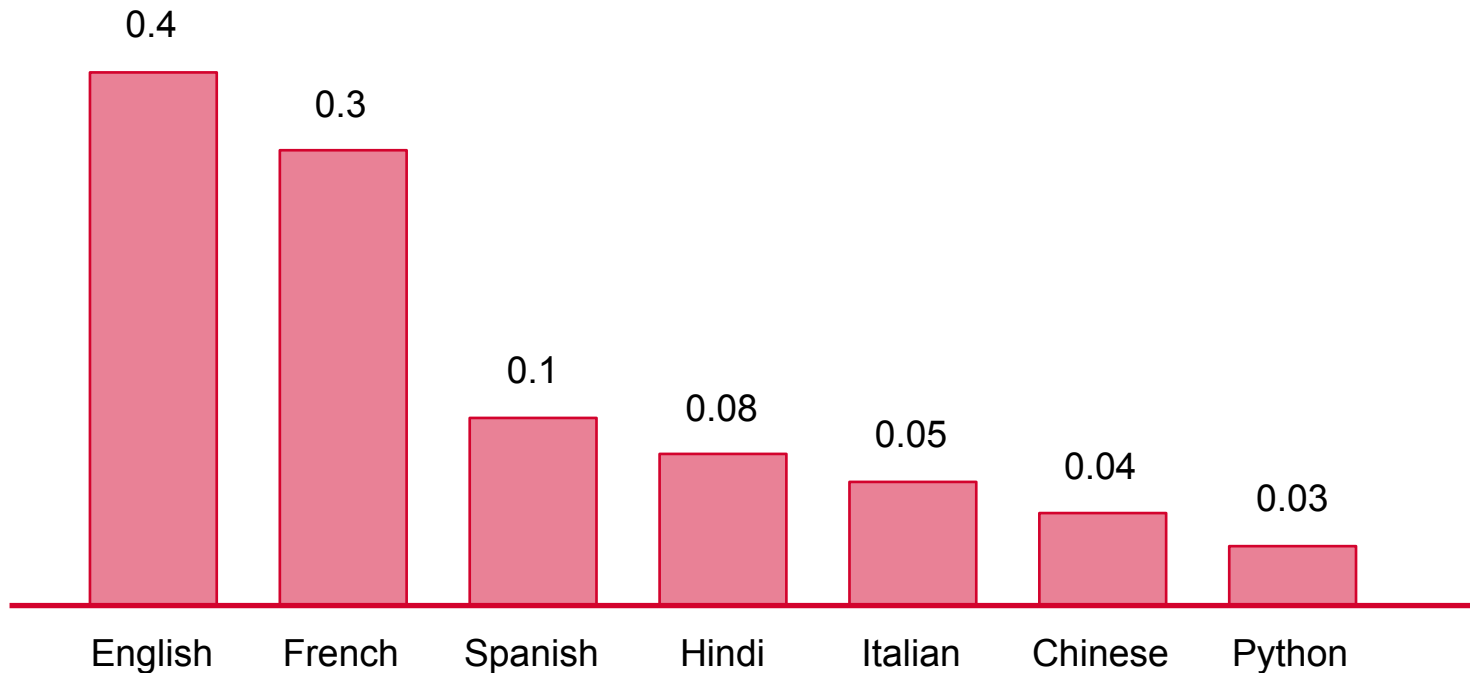




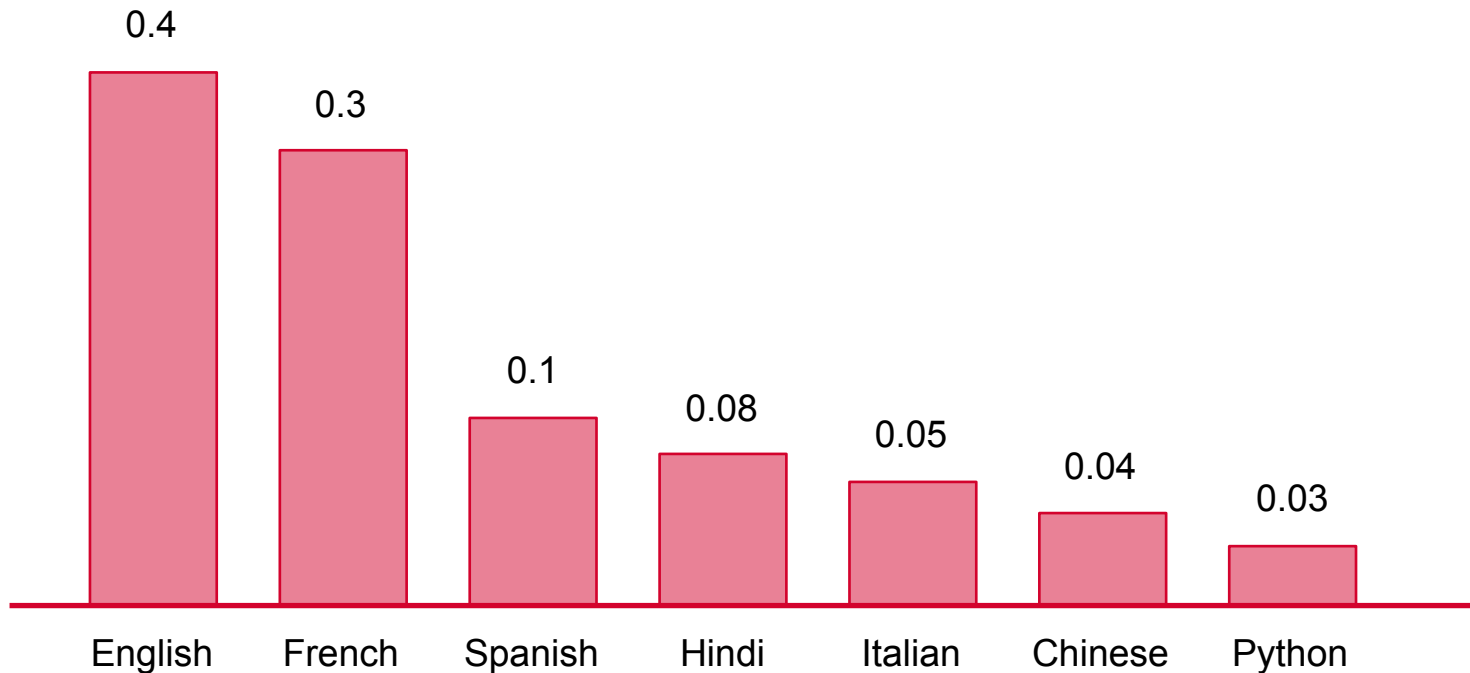
Top-p (Nucleus Sampling)

Top-p sampling, also known as nucleus sampling, works by selecting the **smallest set of top candidate words whose cumulative probability exceeds a given threshold p .**

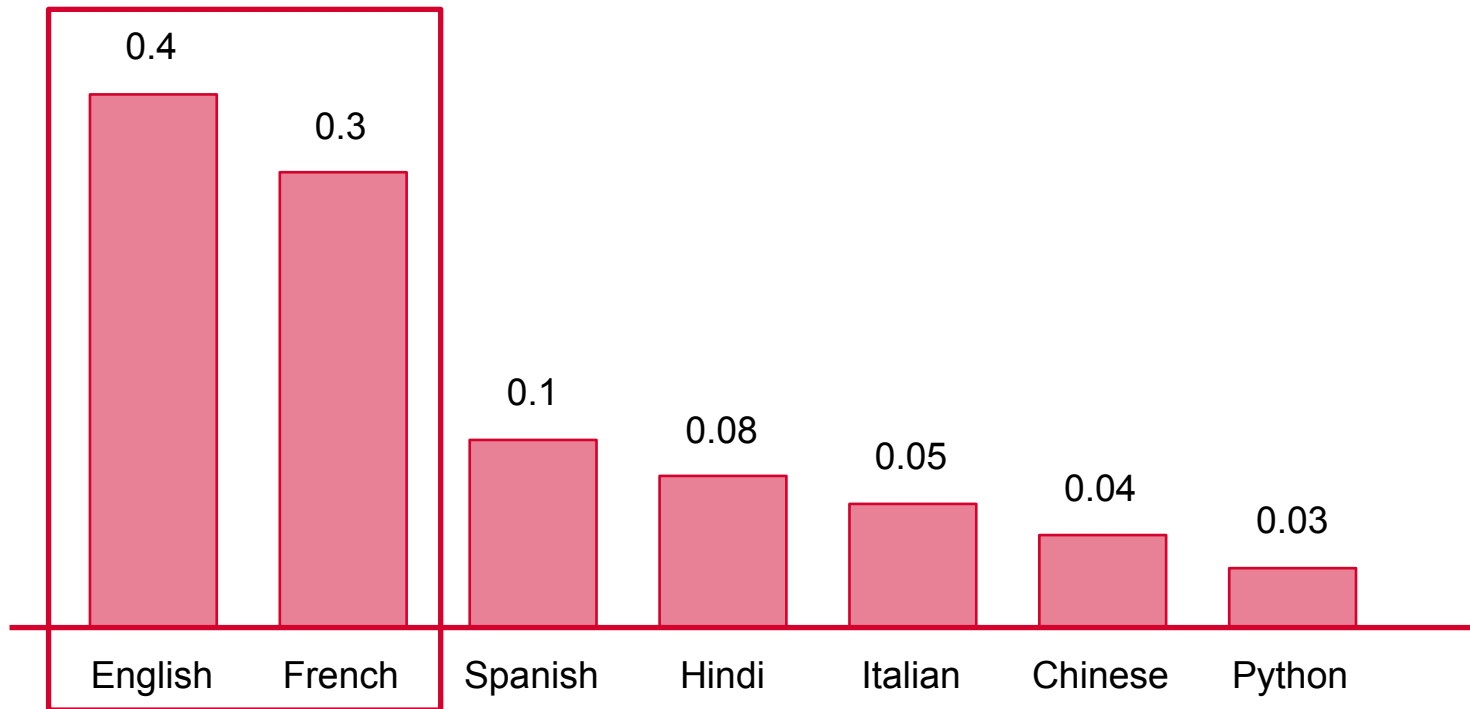
Probabilities of Possible Next Words



Sort All Words in Order of Probability

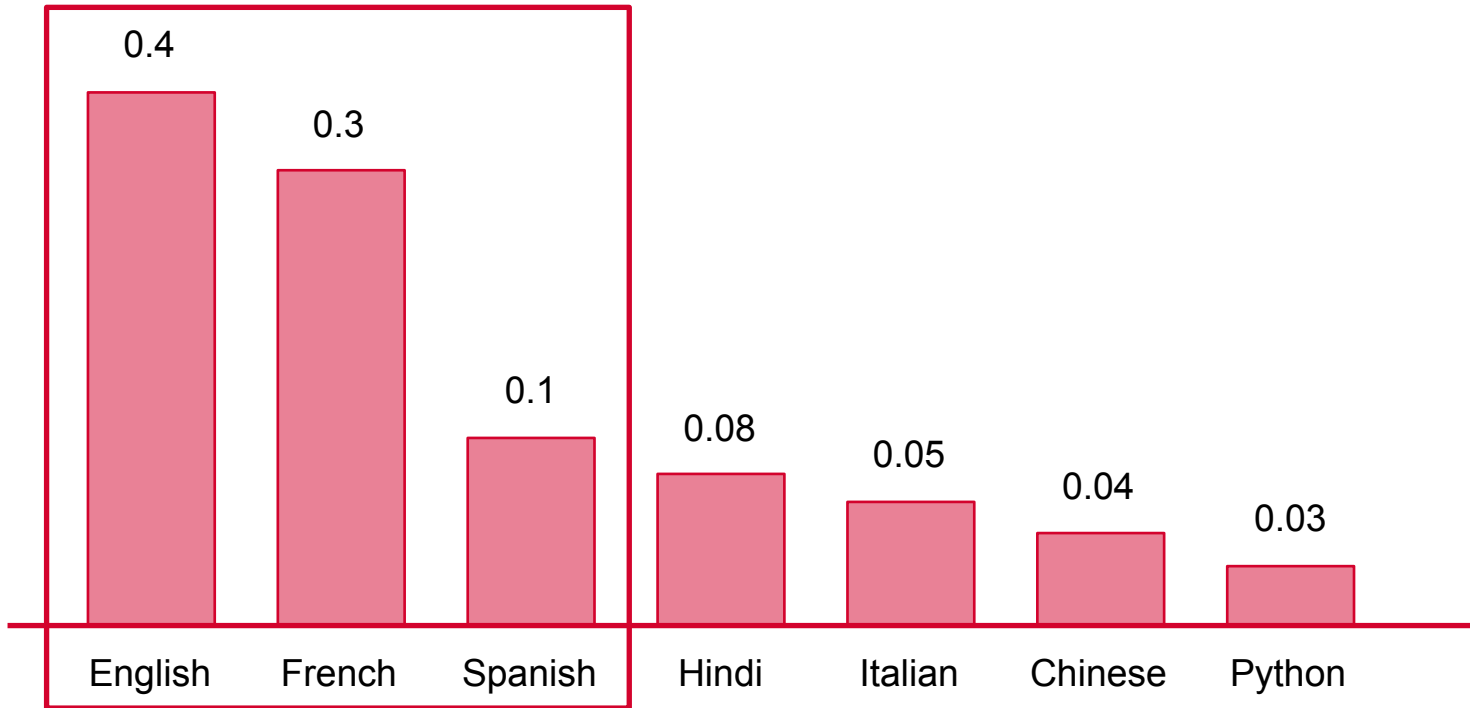


Top-p of 0.5



The next word selected will only choose between the smallest subset of words that exceeds the cumulative probability threshold

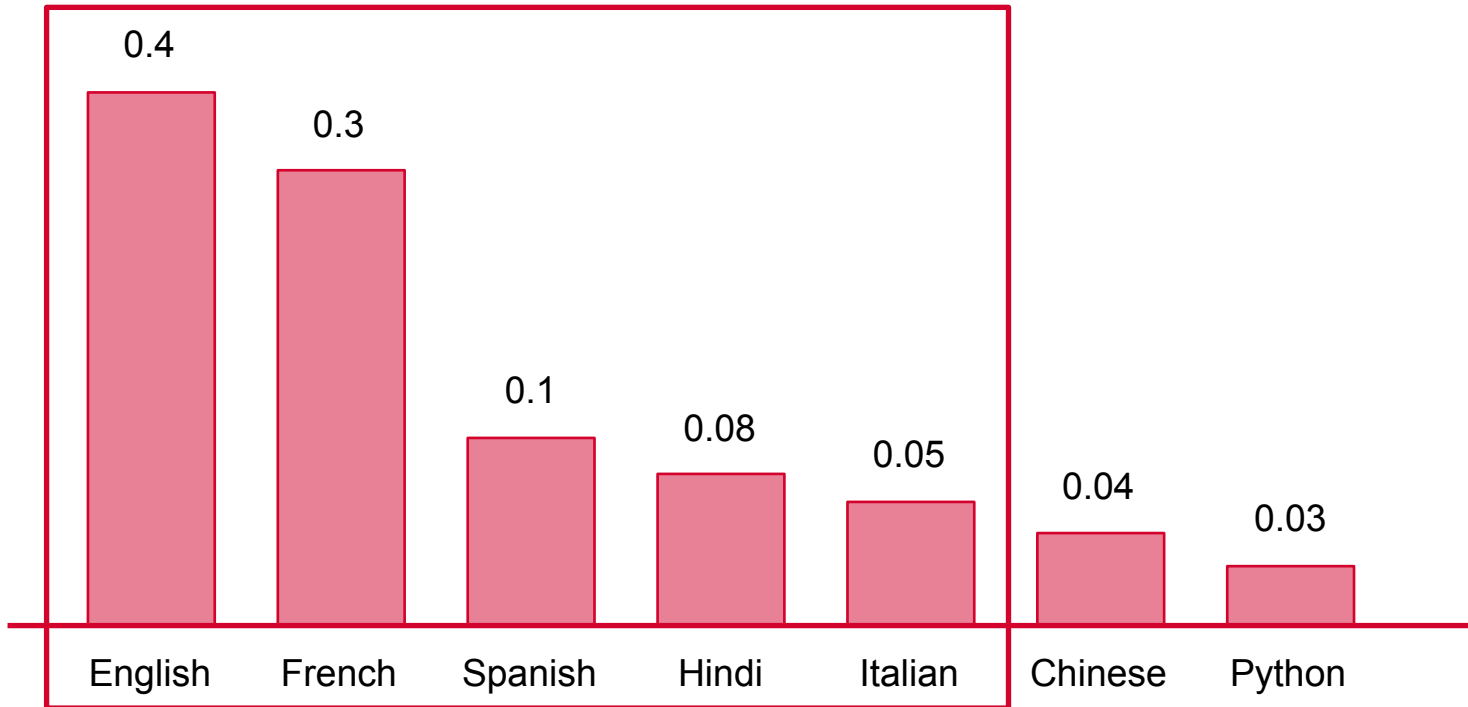
Top-p of 0.7



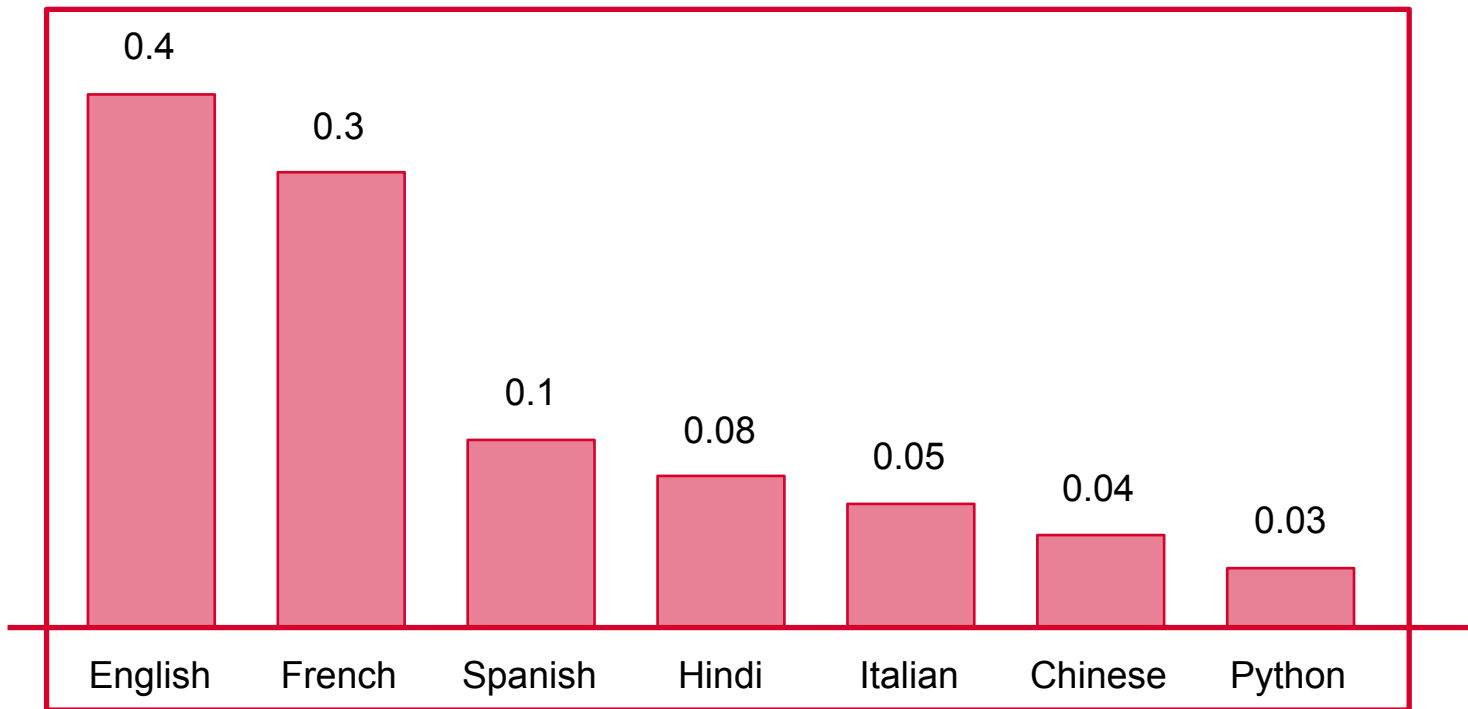
Top-p of 0.9



More words to choose from, more diversity in the output



Top-p of 1



Choose from among all possible words, greatest possible diversity



Top-k Sampling

- Words to choose from limited to the K most likely words at each step
- Value is any positive integer
- Smaller values will pick more likely words, so less diverse output
- Larger values will pick less likely words, more diverse output

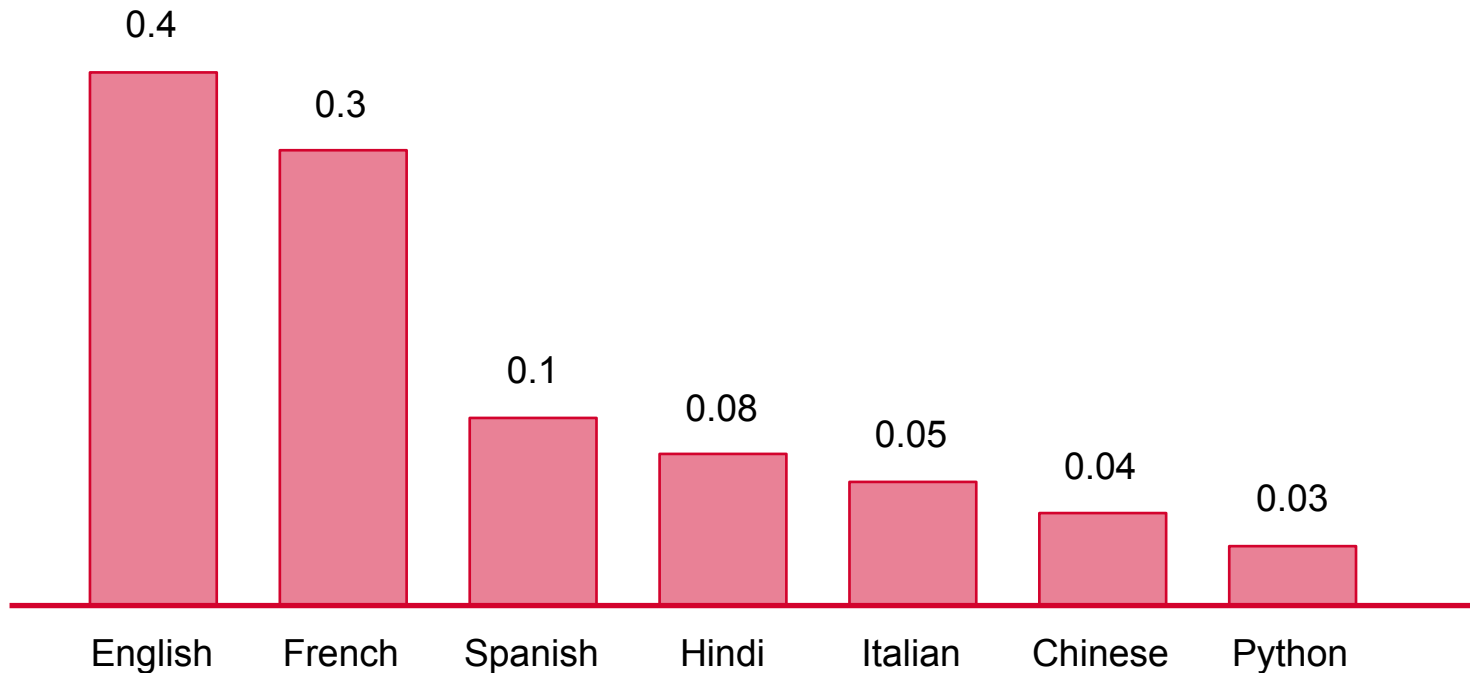





Top-k Sampling

Top-K sampling limits the next-word selection to the **K most likely words** predicted by the model at each generation step.

Sort All Words in Order of Probability




Top-k of 3



The next word selected will only choose between the top 3 words in terms of probability

Top-k of 6



The next word selected will only choose between the top 6 words in terms of probability